

CPE 310

Assembler Project

Final Testing & Robustness Improvements

4/9/2025

Emily Allegeyer, Nicholas Bennet, Ryan Busch, Alexander White

Introduction

While preparing for its Titan-9 launch, our company ByteForge's testing division came across an important issue. There were several bit flips present in the firmware that impacted functionality. We were tasked with identifying and rectifying these bit flips to ensure the Titan-9 can launch successfully.

Experiment

In order to fix the provided machine code, we first had to make it more easily human-readable. This meant parsing the original machine code (shown in *Figure 1*) and noting what each instruction is meant to do via the opcode (first 6 bits) and in some cases the funct value (last 6 bits). From there we split the individual instructions into their smaller parts, such as rs, rt, and immediate. The results of doing this for each instruction are shown in *Figure 2*.

After this, figuring out which bits were flipped was a relatively simple task. For instance, it was clear on the 11th line that the funct field was incorrect because it implied an operation we had no support for. It was observations like this we had to make in order to fix the machine code fully and create the intended product.

```
00000001001010100101000000100000 -> Rd field
00100001000010110000000000000101
00000001001010100110000000100100
00110001110011010000000000001111 -> Rs field
0001000110100000000000000110100
00010101011011010000000000110100
00000001110010000000000000011010
00000000000000100111100000010000 -> Rt field
0000000000000000110000000010010
0000000100101010000000000011000
00000001011011001100100000100001 -> Funct field
00110111001100000000000011111111
00000001000010011000100000101010
00101001000100100000000000001010
00000001001010101001100000100010
00111100000101110100000000000000
10001110111101000000000000000000
10001110111101010000000000000100 -> opcode field
00111100000101100001001000110100
```

Figure 1: Original machine code

ADD	000000 01001 01010 01010 00000 100000	->	Rd field
ADDI	001000 01000 01011 0000000000000101		
AND	000000 01001 01010 01100 00000 100100		
ANDI	001100 01110 01101 0000000000001111	->	Rs field
BEQ	000100 01101 00000 0000000000110100		
BNE	000101 01011 01101 0000000000110100		
DIV	000000 01110 01000 00000 00000 011010		
MFHI	000000 00000 00010 01111 00000 010000	->	Rt field
MFLO	000000 00000 00000 11000 00000 010010		
MULT	000000 01001 01010 00000 00000 011000		
ADDU	000000 01011 01100 11001 00000 100001	->	Funct field
ORI	001101 11001 10000 0000000011111111		
SLT	000000 01000 01001 10001 00000 101010		
SLTI	001010 01000 10010 0000000000001010		
SUB	000000 01001 01010 10011 00000 100010		
LUI	001111 00000 10111 0100000000000000		
LW	100011 10111 10100 0000000000000000		
LW	100011 10111 10101 0000000000000100	->	opcode field
LUI	001111 00000 10110 0001001000110100		

Figure 2: Parsed, human-readable machine code

ADD	000000 01001 01010 01000 00000 100000	ADD	\$t0, \$t1, \$t2
ADDI	001000 01000 01011 0000000000000101	ADDI	\$t3, \$t0, #0x5
AND	000000 01001 01010 01100 00000 100100	AND	\$t4, \$t1, \$t2
ANDI	001100 01100 01101 0000000000001111	ANDI	\$t5, \$t4, #0xF
BEQ	000100 01101 00000 0000000000110100	BEQ	\$t5, \$zero, #0x34
BNE	000101 01011 01101 0000000000110100	BNE	\$t5, \$t3, #0x34
DIV	000000 01110 01000 00000 00000 011010	DIV	\$t0, \$t6
MFHI	000000 00000 00000 01111 00000 010000	MFHI	\$t7
MFLO	000000 00000 00000 11000 00000 010010	MFLO	\$t8
MULT	000000 01001 01010 00000 00000 011000	MULT	\$t1, \$t2
OR	000000 01011 01100 11001 00000 100101	OR	\$t9, \$t3, \$t4
ORI	001101 11001 10000 0000000011111111	ORI	\$s0, \$t9, #0xFF
SLT	000000 01000 01001 10001 00000 101010	SLT	\$s1, \$t0, \$t1
SLTI	001010 01000 10010 0000000000001010	SLTI	\$s2, \$t0, #0xA
SUB	000000 01001 01010 10011 00000 100010	SUB	\$s3, \$t1, \$t2
LUI	001111 00000 10111 0100000000000000	LUI	\$s7, #0x4000
LW	100011 10111 10100 0000000000000000	LW	\$s4, #0x0(\$s7)
SW	101011 10111 10101 0000000000000100	SW	\$s5, #0x4(\$s7)
LUI	001111 00000 10110 0001001000110100	LUI	\$s6, #0x1234

Figures 3 (left), 4 (right): Fixed machine code (left) and assembly instructions (right)

Results

After parsing each line and looking at the denoted fields where the bit flips happened, we believe we successfully produced the machine code that was intended by ByteForge (*Figure 3*). In the first line, the rd field was the one with the bit flip. The flip resulted in a change from storing the result in register \$t0 to storing it in \$t2. This was obvious both because of the pattern of registers used, and the fact that it \$t0 is a source register used in the next instruction. The next bit flip to occur was in the rs field (source register) of the 4th instruction. This would create an incorrect result if left unfixed, so it needed to be changed. The process for figuring out this bit flip was similar to the previous one but reverse. You could tell which register was meant to be the source because of the destination register of the previous instruction. The 01110 had to be changed to a 01100.

The next bit flip was also easy to identify. Although it likely would not have caused any errors due to being in an unused field, it was still best to change. This bit flip was in the rt field of the 8th instruction, an MFHI instruction. As previously stated, the MFHI instruction does not use the rt field so it was a simple change from 01000 to 00000. After that was probably the hardest to fix. It was easy to identify because it was a bit flip in the funct field that implied an operation there was no support implemented for, an error that could crash any Titan-9 systems if unfixed. Figuring out what to change it to, however, was not as easy. There were two possible things the original funct field could have been, 100100 (AND instruction) or 100101 (OR instruction). Either of them would result in a valid assembly instruction, but based on the pattern of the rest of the instructions it made the most sense to change it to 100101.

The 5th and final bit flip was on line 18 and it was the opcode. The corruption turned it into a lw instruction with opcode 100011, which if unfixed would mess with register values in unintended ways. This was also a very easy instruction to fix because there was only one option that made sense: 101011, the sw operation. This made sense because it was the one we had functionality implemented for, and it was one that used an immediate field. After changing this instruction, we believe we fully fixed the corrupted machine code. The results can be seen in *Figure 3*, and were able to be successfully run through our interpreter to get the results in *Figure 4*. This shows that

at the very least, the instructions are all syntactically correct and when reading them in their assembly format look like they create a valid instruction set.